

# Marmara University

## Faculty of Engineering



### CSE 3033

#### Operating Systems

---

## Project 3

---

**Instructor:** Ali Haydar ÖZER

Date:03.01.2024

	Department	Student Id Number	Name & Surname
1	CSE	150120012	Kadir BAT
2	CSE	150120055	Muhammed Talha KARAGÜL
3	CSE	150121021	Feyzullah ASILLIOĞLU

# CONTENTS

<b>CONTENTS.....</b>	<b>2</b>
<b>Abstract.....</b>	<b>3</b>
<b>Questions:.....</b>	<b>4</b>
1. Which method(s) provide the correct result and why?.....	4
2. Among the method(s) providing the correct result, which method is the fastest?.....	4
3. Does increasing the number of threads always result in a smaller total time?.....	4
4. Are there any differences in user time/system time ratio of the processes as the number of threads increases? What could be the cause of these differences?.....	4
<b>Outputs.....</b>	<b>5</b>
Method 1.....	5
Method 2.....	6
Method 3.....	7
Other Outputs.....	8
<b>To Compile Code:.....</b>	<b>8</b>
<b>To Run Code:.....</b>	<b>8</b>

# Abstract

The provided code implements three different methods for calculating the square root sum of a range of numbers concurrently using multiple threads. The methods differ in how they handle synchronization:

- **Method 1:** Each thread independently updates the global sum without synchronization.
  - **Expected Behavior:** For the same input, we expect the results (sum) to be different when the number of threads changes. This is because, without mutex synchronization, threads do not work concurrently. Consequently, Method 1 is prone to race conditions, and the absence of synchronization may lead to incorrect results.
- **Method 2:** Each thread acquires a mutex before updating the global sum to ensure mutual exclusion.
- **Method 3:** Each thread calculates a local sum, and after the local sum is computed, it acquires the mutex to update the global sum.

## Questions:

### 1. Which method(s) provide the correct result and why?

Method 2 and Method 3 are likely to provide correct results because they use a mutex to synchronize access to the global sum. This ensures that multiple threads do not concurrently update the global variable, avoiding race conditions.

Method 1 is prone to race conditions, and it may produce incorrect results. Multiple threads updating the global variable simultaneously can lead to data corruption.

### 2. Among the method(s) providing the correct result, which method is the fastest?

Method 3 is by far the fastest method. Method 2 is quite slow. Although method 2 and method 3 give the same result, there is a speed difference between them. Method 1 is faster than method 2 but gives wrong results.

### 3. Does increasing the number of threads always result in a smaller total time?

Not necessarily. The speedup achieved by increasing the number of threads depends on the nature of the workload and the available CPU cores.

If the workload is not parallelizable or if the number of threads exceeds the number of available CPU cores, increasing the number of threads may not result in a smaller total time.

### 4. Are there any differences in user time/system time ratio of the processes as the number of threads increases? What could be the cause of these differences?

As the number of threads increases, there is an increase in both user time, indicating more time spent in user-level code, and system time, signifying an escalation in kernel-level operations.

Increasing the number of threads may lead to more CPU contention, affecting the ratio. Mutex operations and synchronization mechanisms can contribute to kernel-level activities, impacting the ratio.

# Outputs

We ran the program on 3 different computers. One of them has Linux and the other has a Mac operating system. On the 3rd computer, we installed Linux with a virtual machine and tried it. We got the results the fastest on Mac, so we filled the tables with the data we got from Mac. On the other two computers we had to wait very long times for methods 1 and 2, so we put some of the data we ran on Linux at the bottom of the report.

## Method 1

```
● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 1 1
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 1 1 34.68s user 0.02s system 99% cpu 34.763 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 2 1
Sum: 2.021011e+16
./project3.out 880130203012 922823372203 2 1 136.61s user 0.04s system 199% cpu 1:08.50 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 4 1
Sum: 1.031340e+16
./project3.out 880130203012 922823372203 4 1 143.02s user 0.07s system 396% cpu 36.085 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 8 1
Sum: 5.355896e+15
./project3.out 880130203012 922823372203 8 1 180.96s user 0.22s system 788% cpu 22.971 total
```

Method 1	Sum	User Time	System Time	Total Time	Cpu usage
c = 1	4.053457e+16	34.68 s	0.02 s	34.763 s	%99
c = 2	2.021011e+16	136.61 s	0.04 s	68.50 s	%199
c = 4	1.031340e+16	143.02 s	0.07 s	36.085 s	%396
c = 8	5.455896e+15	180.96 s	0.22 s	22.971 s	%788

In Method 1, the results were different and inaccurate due to the lack of mutex. Also, due to the lack of mutex, not much time was spent in the kernel. Therefore the system time is very short. As expected, CPU usage and user time increased as the number of threads increased. When  $c = 1$ , the total time of the program took 34 seconds. When  $c = 8$ , it decreased to 23 seconds.

## Method 2

```

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 2 2
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 2 2 266.81s user 275.77s system 158% cpu 5:42.80 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 4 2
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 4 2 406.65s user 1005.97s system 297% cpu 7:55.62 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 8 2
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 8 2 400.96s user 2306.09s system 485% cpu 9:17.26 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 16 2
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 16 2 612.51s user 6181.05s system 905% cpu 12:30.00 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 32 2
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 32 2 628.86s user 6091.46s system 872% cpu 12:50.31 total

```

Method 2	Sum	User Time	System Time	Total Time	Cpu usage
c = 1	4.053457+e16	198.35 s	0.05 s	198.56 s	%99
c = 2	4.053457+e16	266.81 s	275.77 s	342.80 s	%158
c = 4	4.053457+e16	406.65 s	1005.97 s	475.62 s	%297
c = 8	4.053457+e16	400.96 s	2306.09 s	557.62 s	%485
c = 16	4.053457+e16	612.51 s	6181.05 s	750 s	%905
c = 32	4.053457+e16	628.86 s	6091.46 s	770.31 s	%872

In method 2, compared to method 1, we obtained accurate results in all inputs. We expected user time and system time to increase because we use mutex in this method. We did not expect the total time to increase, but as the number of threads increased, the program started to run slower. Total CPU usage would be very high with the increase in the number of threads as expected. We tried this method on a Linux computer with a not-very-powerful processor. The program ran for a very long time. That's why we took the data with Mac.

## Method 3

```

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 1 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 1 3 35.49s user 0.02s system 99% cpu 35.630 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 2 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 2 3 35.50s user 0.02s system 199% cpu 17.851 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 4 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 4 3 71.10s user 0.14s system 351% cpu 20.252 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 8 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 8 3 95.31s user 0.13s system 784% cpu 12.165 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 16 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 16 3 65.38s user 0.14s system 930% cpu 7.043 total

● karagul@Muhammed-MacBook-Pro CSE3033_Project3_2023 % time ./project3.out 880130203012 922823372203 32 3
Sum: 4.053457e+16
./project3.out 880130203012 922823372203 32 3 69.68s user 0.12s system 947% cpu 7.366 total

```

Method 3	Sum	User Time	System Time	Total Time	Cpu usage
c = 1	4.053457+e16	35.49 s	0.02 s	35.63 s	%99
c = 2	4.053457+e16	35.50 s	0.02 s	17.851 s	%199
c = 4	4.053457+e16	71.10 s	0.14 s	20.252 s	%351
c = 8	4.053457+e16	95.31 s	0.13 s	12.165 s	%784
c = 16	4.053457+e16	65.38 s	0.14 s	7.043 s	%930
c = 32	4.053457+e16	69.68 s	0.12 s	7.366 s	%947

Method 3 was the fastest and most accurate method. It worked very fast compared to Method 2. It gave more accurate answers than Method 1. Therefore, method 3 was the best method among these 3 methods. As the number of threads increased, user time increased. This is a normal result. System time was quite short despite the increase in the number of threads. For this reason, the total time of the program was also quite short. We expected the program to speed up as the number of threads increased, and it did. The program that took 35 seconds when  $c = 1$  took 7 seconds when  $c = 32$ . We learned how useful threads are.

## Other Outputs

```
• karagul@karagul-virtual-machine:~/Desktop/GitHub/CSE3033_Project3_2023$ time ./project3.out 880130203012 922823372203 1 1
Sum: 4.053457e+16

real    2m34,072s
user    2m33,846s
sys     0m0,182s

• karagul@karagul-virtual-machine:~/Desktop/GitHub/CSE3033_Project3_2023$ time ./project3.out 880130203012 922823372203 2 1
Sum: 2.358994e+16

real    5m17,873s
user    10m15,646s
sys     0m10,325s

sys     0m0,003s
• kadir@kadir-Rev-B:~/Desktop/CSE3033_Project3_2023$ time ./multithreaded_sum 880130203012 922823372203 1 2
Sum: 4.053457e+16

real    11m56,591s
user    11m56,133s
sys     0m0,141s
• kadir@kadir-Rev-B:~/Desktop/CSE3033_Project3_2023$

• kadir@kadir-Rev-B:~/Desktop/CSE3033_Project3_2023$ time ./multithreaded_sum 880130203012 922823372203 2 2
Sum: 4.053457e+16

real    30m58,072s
user    39m22,639s
sys     21m51,817s
• kadir@kadir-Rev-B:~/Desktop/CSE3033_Project3_2023$
```

These are some outputs that ran in linux. The program took too long to run because the Linux computer had a bad processor.

## To Compile Code:

➤ gcc project3.c -o project3.out -lm -pthread

## To Run Code:

➤ time ./project3.out `**<startOfTheRange>**` `**<endOfTheRange>**` `**<threadNumber>**`  
`**<methodNumber>**`